Supplemental materials for "What's past is present: Reminders of past choices bias decisions for reward in humans."

Aaron M. Bornstein [1*], Mel W. Khaw [3], Daphna Shohamy [4] & Nathaniel D. Daw [1,2]

1. Neuroscience Institute and 2. Department of Psychology, Princeton University, Princeton, NJ, USA.

3. Department of Economics, Columbia University, New York, NY, USA.

4. Department of Psychology; Zuckerman Mind, Brain, Behavior Institute;

and Kavli Center for Brain Science, Columbia University, New York, NY, USA.

* To whom correspondence should be addressed: aaronmb@princeton.edu.

Supplemental results

Behavior

This section describes variants of the models presented in the main text.

Regression weights consistent on simulated model

We show that the regression results follow from the episodic sampling model. To simulate the model, we generated 50 populations of 20 simulated subjects, each of whom ran a unique instantiation of the task (with different payoff timeseries and outcomes), and fit the regression model to each population. Simulated subjects drew one sample before each decision, used the mean choice parameters as fit to the human population, and gave, on average, accurate responses to memory probes at the same rate as did real subjects. Figure S1 shows the average regression weights, across these populations, for each variable of interest.

Alternative forms of choice noise

One potential explanation for the superior fit of the sampling model is that it simply captures additional stochasticity in subjects choices, over and above that captured by the standard softmax choice function. For instance, subjects could, with some probability ϵ, select the highest valued option, rather than selecting based on the difference in value between the two options ([1]; Equation S1).

$$p_t(a = A_i) = \epsilon(Q_{t-1}^{TD}(A_i) = max(Q_{t-1}^{TD}(\cdot))) + (1 - \epsilon)(\frac{e^{\beta^c I_t^c + \beta^{TD} Q_t^{TD}(A)}}{\sum_j e^{\beta^c I_t^c + \beta^{TD} Q_t^{TD}(a_j)}}) \tag{S1}$$

However, model comparison did not provide evidence in favor of this "ϵ-greedy" approach. In Experiment 1 the Sampler model was favored for 15/20 subjects, by a mean Bayes Factor of 3.0042 (SEM 1.8137, exceedance probability > 0.99), while in Experiment 2 the Sampler was favored for 19/21 subjects, by a mean Bayes Factor of 6.734 (SEM 1.7078, exceedance probability > 0.99).

Adding evoked trials to the TD model

We augmented the TD model to incorporate rewards from bandit trials evoked by valid memory probes. Specifically, we added an additional parameter, α^{evoked}, for a new, augmented TD update applied to rewards $r_{i_{evoked}}$ evoked during memory probe trials (Equation S2). This parameter allowed the weight given to evoked
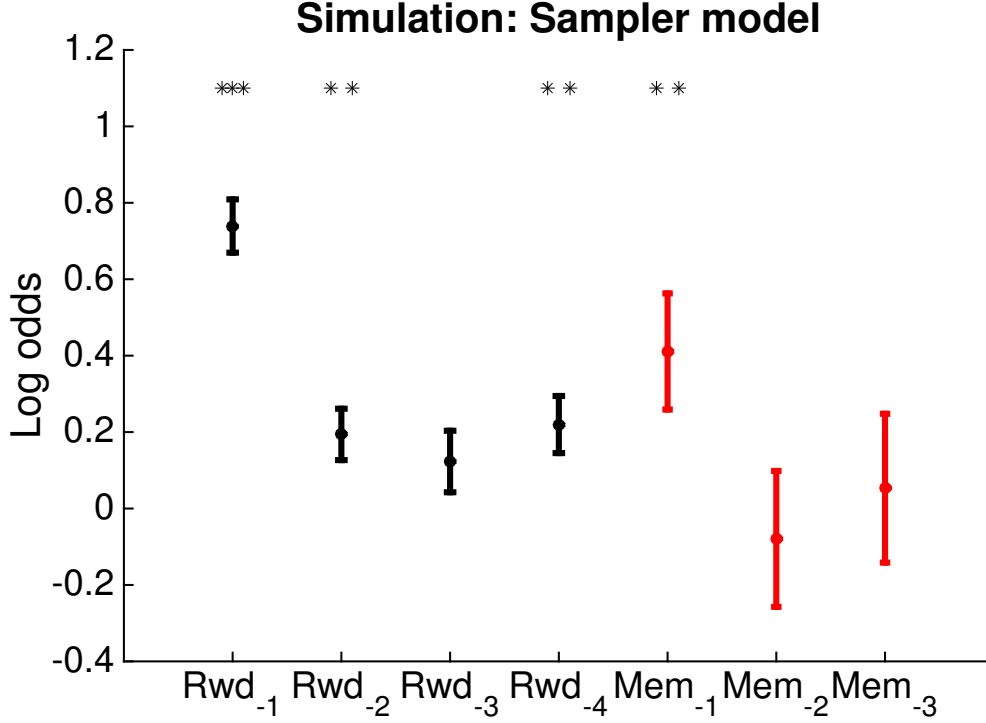
Figure S1: **Regression analysis on simulated subjects.** Average regression results for 50 populations of 20 simulated subjects each. Subjects were simulated using the Sampler model at the parameters fit to the study population. Error bars are SEM across the population means.

bandit outcomes to vary, reflecting the fact that the sampling mechanism may itself be stochastic in nature — not every probe trial will successfully trigger a recall of the associated context, even those on which participants exhibit correct recognition memory.

$$Q_t^{TD}(a) = Q_{t-1}^{TD}(a) + \alpha^{evoked}(r_{i_{evoked}} - Q_{t-1}^{TD}(a)) \tag{S2}$$

Table S1 expands the comparison from the main text to include this *TD-evoked* model. After correcting for the additional parameter, the model was a slightly better fit to participants behavior than the plain TD model. It was not, however, a superior fit to the Sampler model.

Model	α	α^{evoked}	β	β^c	log Bayes
TD	0.5552 (0.0862)	-	1.7551 (0.6845)	-0.0930 (0.2354)	6.9182 (1.3227)
TD-evoked	0.5269 (0.0842)	0.2135 (0.0545)	2.3351 (0.7223)	-0.0962 (0.2381)	5.9167 (1.2677)
Sampler	0.5393 (0.0583)	0.4386 (0.0990)	2.2869 (0.4943)	0.5855 (0.3215)	-

Table S1: **Fit model parameters for Experiment 2, expanded to include the TD-evoked model** The parameters shown are the mean (SEM) across subjects. The final column shows the mean (SEM) of the log Bayes Factor versus the Sampler model (smaller is better).

Hybrid Sampler and TD model

The evidence across several studies shows that multiple valuation systems contribute to choice, either simultaneously or across time [2–4]. The current study provides evidence that one component of this valuation architecture involves evaluating samples from episodic memory. (Of course, we do not know to what extent this influence is coextensive with, e.g. model-based learning as otherwise defined.)

To test the possibility that such a "hybrid" model could account for choices in this task, we implemented a hybrid model combining the TD and episodic sampling models and tested it on Experiment 2. The model has six parameters: a learning rate for the TD component α^{TD}, a decay rate for the Sampler component α^{Sample}, a choice stickiness term β^c, a softmax temperature for the TD value β^{TD}, a softmax temperature for the β^{Sample}, and a decay rate for evoked trials α^{evoked}. Choice probability is computed using Q-values derived from each model, as specified in the main text, and taken over all possible combinations of samples (following Equation 5 in the main text).

Parameter estimates from this model show that that most of the weight is on the sampler model (in the sense that it has a much higher softmax temperature, i.e. its values have a larger effect on choice – indeed, the weight on the TD model is not reliably different from zero). Accordingly, the addition of a TD component to the sampler model was not robustly justified in light of the additional free parameters. On average, across subjects, log Bayes Factors were mildly in favor of the Sampler model (mean 0.9700, SEM 1.1015, exceedance probability 0.8999), and the Sampler was a better fit for 13 out of 21 subjects individually. The fit parameters and model comparison results are shown in Table S2.

α^{TD}	α^{Sample}	β^c	β^{TD}	β^{Sample}	α^{evoked}	log Bayes
0.4275 (0.0653)	0.5670 (0.0521)	0.5281 (0.2815)	0.0580 (0.4654)	2.0910 (0.5187)	0.6005 (0.0871)	0.9700 (1.1015)

Table S2: **Fit model parameters for the Hybrid model.** The parameters shown are the mean (SEM) across subjects. The final column shows the mean (SEM) of the log Bayes Factor versus the Sampler model.

In sum, after accounting for the additional free parameters, this hybrid model was not a clearly superior fit to behavior than the episodic sampling model taken alone. However, we think that–in an experiment designed to distinguish these two possibilities–a more sophisticated architecture (perhaps employing a common cached value representation as in DYNA [5]) could possibly prove a superior explanation of behavior.

Distinguishing models on simulated data

To demonstrate that these models are, in fact, distinguishable, we simulated the models each running 1,000 instantiations of each experiment, each instance with separately initialized payoff and outcome timeseries. Each model simulation was run using parameters as fit to one randomly drawn subject from the respective Experiment. The Sampler model drew one sample before each choice. For Experiment 2, simulated subjects responded to memory probes correctly the same percentage as did our real subjects. We then fit both models to each population of 1,000 simulated subjects. The result of these fits is shown in Table S3.

For Experiment 1, subjects simulated using the TD model, 81.5% were best-fit by the TD model, by an average log Bayes factor of 12.8726 (SEM 0.8499). For Experiment 1, subjects simulated using the Sampler model, 89.7% were best-fit by the Sampler model, by an average log Bayes factor of 8.4295 (SEM 0.7734).

For Experiment 2, subjects simulated using the TD model, 88.7% were best-fit by the TD model, by an average log Bayes factor of 5.9292 (SEM 1.5065). For Experiment 2, subjects simulated using the Sampler model, 91.3% were best-fit by the Sampler model, by an average log Bayes factor of 3.5762 (SEM 0.6757).

In both datasets, the corresponding simulated model was a superior fit, for the bulk of the population and on average at the individual level.

Neuroimaging

Given the Sampler's superior fit to behavior, we used the neuroimaging data collected alongside Experiment 1 [6] to ask whether the expectation and learning variables predicted by the sampling model could provide
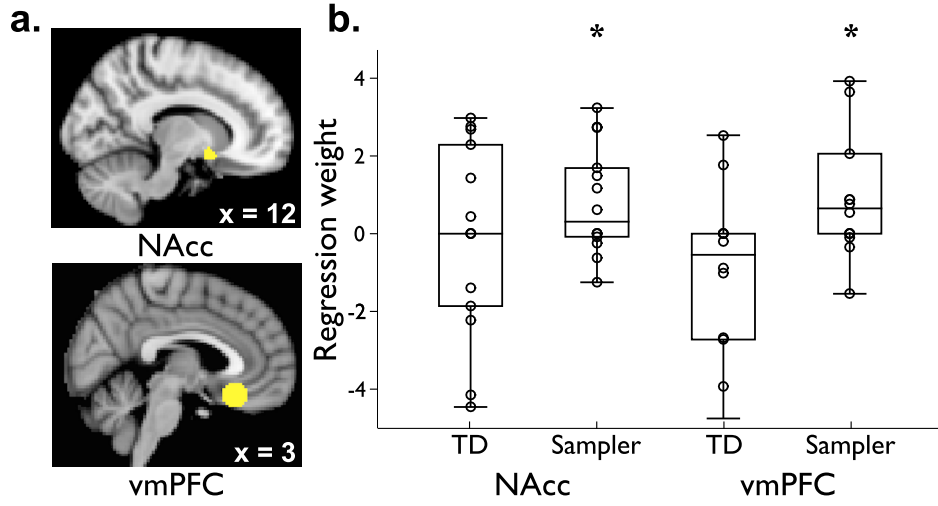
Figure S2: Sampling fit to neural decision variables a. Regions of interest. We isolated voxels of interest that corresponded to previous reports of the neural substrates for the decision variables analyzed here: Chosen Value (CV), in ventromedial prefrontal cortex (vmPFC), and Reward Prediction Error (RPE), in the nucleus accumbens (NAcc). b. Simultaneous regression. Candidate timeseries for each decision variable were generated according to each of the two models, and entered into a simultaneous regression against the BOLD timeseries extracted from the relevant ROI. Each plotted point represents the regression coefficient for the respective model-timeseries pair; box plots display the mean and interquartile range (* $p < 0.05$). Both regressions support the hypothesis that the Sampler model underlies neural signals (NACC-RPE: $t_{13} = 2.2134$, $p = 0.0454$; vmPFC-CV: $t_{13} = 2.2604$, $p = 0.0416$).

4

Task	Simulated model	Fraction best-fit	log Bayes
Expt 1	TD	0.815	12.8726 $_{(0.8499)}$
	Sampler	0.897	8.4295 $_{(0.7734)}$
Expt 2	TD	0.887	5.9292 $_{(1.5065)}$
	Sampler	0.910	3.5762 $_{(0.6757)}$

Table S3: **Confusion matrix for Sampler and TD models.** For each experiment and each model, we simulated 1,000 participants using the given model as ground-truth. Each individual simulated participant used a set of parameters selected at random from the parameters fit to human participants. Both models were then fit to each simulated participant's choices. Shown are the fraction of simulated participants best-fit by the ground-truth model and the mean (SEM) log Bayes factor in favor of that model.

a better explanation of BOLD signal than did the corresponding variables extracted from a TD model. Specifically, we tested whether the well-studied neural correlates of key decision variables–Chosen Value (CV) and Reward Prediction Error (RPE)–were better predicted by the sampling model than by TD. We first identified regions of interest (ROIs) encompassing areas previously shown to reflect this activity: ventromedial prefrontal cortex / medial orbitofrontal cortex (hereafter: vmPFC) for chosen value, and Nucleus Accumbens (NAcc) for RPE (Figure S2a). Extracting the timeseries of activity within these ROIs, we next performed a simultaneous regression containing the timeseries of variables predicted by both models, along with several regressors of no interest. Comparing the distribution of resulting per-participant regression weights against zero using a two-tailed, one-sample t-test, we evaluated whether each model was a significant predictor of the target BOLD timeseries.

The regressors were slightly, but reliably, correlated between models (for RPE: mean $R = 0.1093$, $p = 0.0215$; for CV: mean $R = 0.2701$, $p = 0.0002$). To test the exclusive contribution of the Sampler-derived predictor variables, we orthogonalizd the RPE and CV timeseries as generated using the Sampler against their TD counterparts, and entered each set of TD and Sampler predictors into a simultaneous regression on the BOLD timeseries.

In both cases, the predictions of the Sampler model captured additional variance in the BOLD timeseries that was not modeled by TD (Figure S2b). Across participants, regression on the NAcc timeseries revealed significant contribution of the RPE variable as generated by the Sampler model ($t_{13} = 2.2134$, $p = 0.0454$), but not TD ($t_{13} = -0.1614$, $p = 0.8742$). Similarly, the Chosen Value regressor generated by the Sampler model was a significant predictor of the vmPFC timeseries ($t_{13} = 2.2604$, $p = 0.0416$), while the TD version was not ($t_{13} = -1.0835$, $p = 0.2983$).

References

[1] Hanan Shteingart, Tal Neiman, and Yonatan Loewenstein. The Role of First Impression in Operant Learning. *Journal of Experimental Psychology: General*, 142(2):476–488, aug 2012. ISSN 1939-2222. doi: 10.1037/a0029550. URL http://www.ncbi.nlm.nih.gov/pubmed/22924882.

[2] Nathaniel D Daw, Yael Niv, and Peter Dayan. Uncertainty-based competition between prefrontal and dorsolateral striatal systems for behavioral control. *Nature Neuroscience*, 8(12):1704–1711, 2005.

[3] Nathaniel D Daw, Samuel J Gershman, Ben Seymour, Peter Dayan, and J Raymond. Model-based influences on humans' choices and striatal prediction errors. *Neuron*, 69(6):1204–1215, 2011. doi: 10.1016/j.neuron.2011.02.027.Model-based.

[4] Aaron M. Bornstein and Nathaniel D. Daw. Cortical and Hippocampal Correlates of Deliberation During Model-Based Decisions for Rewards in Humans. *PLoS Computational Biology*, 9(12):e1003387, dec 2013. ISSN 1553-7358. doi: 10.1371/journal.pcbi.1003387. URL http://dx.plos.org/10.1371/journal.pcbi.1003387.

[5] Richard S. Sutton. Dyna, an integrated architecture for learning, planning, and reacting. *ACM SIGART Bulletin*, 2(4):160–163, jul 1991. ISSN 01635719. doi: 10.1145/122344.122377. URL `http://portal.acm.org/citation.cfm?doid=122344.122377`.

[6] N D Daw, J P O'Doherty, P Dayan, B Seymour, and R J Dolan. Cortical substrates for exploratory decisions in humans. *Nature*, 441:876–879, 2006.